

# Assignment 2: Genetic Programming for Classification and Transfer learning

OMPHEMETSE SENNA

ARTIFICIAL INTELLIGENCE

## 1. Abstract

### 1.1. Purpose

The objective of this study is to investigate the application of transfer learning to a target problem dataset, with the aim of assessing the extent to which knowledge acquired from a previous dataset(diabetes) can be effectively transferred and utilized to improve the performance on the new related data or target. By leveraging pre-existing knowledge and models, it is hypothesized that there will be an increase in accuracy and efficiency in handling the new dataset. The experimental design includes the assessment of various transfer learning techniques to ascertain the most effective approach in enhancing model generalization and performance.

## 2. Introduction

**D**iabetes is a chronic metabolic pathology characterized by the fact that the body cannot regulate sugar levels for a long period due to insufficient or weak synthesis of insulin . Increased blood sugar levels provoke the development of various health problems from kidney damage to cardiovascular diseases. Diabetes is genetically inherited. Furthermore, various forms of diabetes can also be caused by obesity, unbalanced nutrition, and the victim's age, especially if the person is elderly . All of these factors are crucial for understanding and preventing diabetes. Lifestyle changes allowing reducing blood sugar levels achievement. They include regular physical activity, weight loss, and a balanced diet.

Finally, blood sugar regulation requires someone to take medications. Blood sugar monitoring must occur regularly to ensure one can make informed judgments about medicine and exercise and eating. Public instruction campaigns and laws and other preventative measures, such as safeguards, are critical. Governments need to spend a lot of money on health care as a result of these diseases. Diabetes affects millions of people around the world. To address this unpleasant expansion, all people, the medical staff, and the legislature must coordinate their efforts.

One of the commonly used techniques in machine learning is transfer learning, which refers to the ability to apply the knowledge gained in a given task to another related task and boost the learning performance. The paper examines a novel framework for transfer learning with respect to genetic programming for classification employing trees, including decision trees.

## 3. Data set, data pre-processing, feature extraction

### 3.1. Attributes

The original dataset or the source data set incorporates a set of health-related factors, namely, pregnancies, glucose, blood

pressure, skin thickness, insulin, BMI, diabetes, age, and pedigree, with which the presence or absence of the disease has been predicted in binary classes.

In contrast, the target dataset also includes pregnancies, glucose, blood pressure, skin thickness, insulin, BMI, diabetes and age, with the target class of Outcome also represented as binary to indicate its presence or absence.

The 'Outcome' class, represented in the source dataset, is the target variable that expresses the presence or absence of diabetes. The 'Outcome' class is binary, as 1 represent the presence of diabetes, while 0 represents the absence. Due to the class's ultimate decision about the prediction, many efforts were taken to pre-process the data and balance the class, which enabled the model to use all instances for more effective and generalizable predictions.

### 3.2. Data Cleaning

Based on the data analysis, there are no missing values in the diabetes dataset, which is our source data used for the study and also our target after gaining insight. Overall, the well-balanced distribution of data points may be deduced. In this regard, means of imputation such as mean, median, and mode replacements or more advanced methods, such as KNN imputation, might be utilised in future analyses with missing data. Therefore, these means of imputation are currently not applicable due to the complete absence of missing data, which is an indicator of the robustness and sparseness of the current dataset.

The dataset was subjected to several data cleaning methods to carefully inspect and analyse it. These methods included the identification of missing values, duplicates, and disparate data formats in the dataset. More specifically, the dataset was tested for missing values. It was found that there are no missing values in the dataset, which means that all the values in the dataset were complete and accurate. Furthermore, there were no duplicates or diverse data formats, implying that the data is useful and error-free for both the source and target problems.

During the data cleaning process, it was discovered that outliers were present in the dataset. This was because, in attempting to remove the identified outliers as shown in figure 1 and 2, the occurrence of null results in the dataset was observed to have increased, thereby complicating the cleaning process. Moreover, it was observed that after the removal of suspected outliers, the same would reappear, and this would mean that the process of removal was not effective.

This process reveals that the removal of outliers is bound to destroy the quality of the data. In conclusion, it was determined that removing these outliers would likely destroy the quality of the data, and therefore, it was decided to

continue with the original dataset as it is, acknowledging the presence of outliers.




Figure 1. Example of a figure




Figure 2. Example of a figure

## 4. Representation

The classifier representation used in our study for the source data (the diabetes dataset), and for the target data (the diabetes dataset) is a decision tree model. A decision tree is graphically depicted as a tree with hierarchical levels in this representation; internal nodes perform features, the branches depict decision rules, and leaf nodes perform class labels. By definition, this representation is understandable and structured. It also ensures that the source data modifies the structure and parameters of the model to improve the decision-making process.

Furthermore, in the context of transfer learning, decision trees are instrumental in transferring knowledge acquired from one domain to another by adapting the decision rules according to the new dataset (target). This approach enables more proficient and effective learning across divergent yet connected domains.

### Decision Tree Representation:

- Each decision node represents a decision based on a feature attribute, where the branch corresponds to a specific condition on that attribute. For example, Glucose  $> 120$ .
- Leaf nodes represent the predicted class labels (e.g., Outcome = 1 or Outcome = 0).

## 5. Functional and terminal data set

**Decision Node:** Represents a decision based on a feature attribute. It has an arity of 2 (binary), where each branch corresponds to a specific condition on the feature attribute.

**Leaf Node:** Represents the class label prediction. It

has an arity of 0 (terminal), where each leaf node corresponds to a specific class label.

### 5.1. Function Set:

**Decision Nodes:** These are internal nodes in the decision tree that split the data based on specific feature conditions.

**Feature Attributes for Source dataset:** Attributes such as Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI, DiabetesPedigreeFunction, and Age from the dataset.

**Feature Attributes for target dataset:** Pregnancies, Glucose, BloodPressure, SkinThickness, Insulin, BMI, DiabetesPedigreeFunction, and Age from the dataset.

**Arity:** Binary (2)

**Value/Range:** Each decision node represents a feature condition, such as "Glucose  $\leq$  threshold" or "BMI  $>$  threshold", where the threshold is a specific value determined during training.

### 5.2. Terminal Set (Decision Tree):

**Leaf Nodes:** These are the terminal nodes in the decision tree that assign a class label to the instances that reach them.

**Terminal Attributes for Source Output.**

**Terminal Attributes for Target Output.**

**Arity:** Unary (1)

**Value/Range:** Each leaf node corresponds to a class label. For binary classification (as in this case), there are two leaf nodes, typically representing class 0 and class 1.

## 6. Fitness function

The fitness function in this research is on classification accuracy; several other classification metrics, including precision, recall, and more, are also generated. The accuracy score is generated by comparing the predicted labels against the actual labels of the test data. The metric provides an indication of the fractional proportion of correct classification cases to the total cases.

Furthermore, precision and recall scores are generated to give further performance information about the classifier other than simple accuracy. Precision significantly measures the total true positive predictions over the sum of positive classifications made by the classifier. Recall, on the other hand, tallies the count of true positive predictions to the sum of actual positive instances in the data.

The primary metric in the evaluation of the efficiency of the individual solutions that genetic programming with and without transfer learning finds for our classification tasks is accuracy, as it also measures the capability of the classifier to predict a positive outcome properly. In such a way, the main feature of the classifier is its sub-models' ability to predict a positive outcome properly.

However, other metrics, which offer a more comprehensive insight into the classifier's quirks under the conditions of

class imbalance, such as precision and recall metrics, are also essential to consider. Hence, while it is vital, accuracy is utilised as the main fitness function; other metrics, such as precision and recall, are also integrated into the evaluation of the performance of the decision trees that were obtained through evolution.

## 7. Selection method

In our study, tournament selection is used as the parent selection method in the genetic programming algorithm with and without transfer learning. Tournament selection is a standard method of sampling individuals in evolutionary algorithms based on their fitness.

Tournament selection involves choosing a fixed number of individuals randomly selected from the population to join a tournament to decide which will be selected for the final sample. This sample will then be put back into the population for further rounds of selection. The participant with the highest fitness is said to have ‘won the tournament’ and is chosen.

This process is repeated for each parent needed for crossover and mutation, ensuring that the fittest individuals have a higher chance of being selected as parents. In the context of the provided code, tournament selection is performed iteratively for each generation to create a new generation of individuals with potentially improved fitness.

## 8. Genetic operators

### Crossover Operator:

#### Description:

Crossover is the genetic operation that merges the genetic information of two parents to produce offspring.

#### Implementation:

For our research, crossover is performed between pairs of parents to give rise to two offspring. The crossover is a stochastic operation through the crossover rate, 0.8, for the parents. Each parent pair will likely be crossed by 80 percent. Specifically, two parents are chosen at random from the population. Next, two random nodes are selected from the decision tree of the parents. Finally, the subtrees rooted at the selected node in the parents are swapped to generate two offspring.

#### Impact:

Crossover combines beneficial genetic material from different parents, potentially producing offspring with improved characteristics compared to their parents.

### Mutation Operator:

#### Description:

Mutation is a genetic operation that alters one or more genes in an individual from the population. For our research, mutation is applied to a single parent, and it modifies a randomly selected node in the decision tree structure.

#### Implementation:

The mutation process is stochastic, characterised by the mutation rate, which is equal to 0.1 in this case. Therefore, for each parent, there is a 10 percent chance of a mutation. If

the random probability is greater than the mutation rate, the node of the decision tree is chosen to mutate.

### There are two types of mutations in our case:

- Left-Child Mutation: With a 0.5 opportunity, the left child of the selected node will be replaced with a randomly chosen node out of the decision tree.
- Right Child Mutation: The right child of the chosen node is substituted with a completely randomised node from the entire tree with a probability of 50 percent.

### Impact:

The mutation process generates a new population that supports the diversity of parts of the search space. Thus, it creates output that assists in exploring new regions.

## 9. Termination of Program:

In this experiment, a parameter was set specifying that the script (code) must reach generation 10. Upon reaching this specified generation, the script will terminate automatically. Subsequently, all generations will be displayed along with their respective evaluations to compare and determine which generations produced better results. This process aims to provide insight into the performance of the program over different generations. The results obtained from this analysis will aid in understanding the evolutionary progression of the program and the effectiveness of its iterations.

## 10. Transfer Learning

### 10.1. What is Transferred:

The pre-trained decision tree models generated from the source problem form the original population for the genetic programming algorithm for the target problem. First, the optimal individual population may be transferred from the source to the target problem in the code execution. This entails conducting tournament selection during genetic programming’s initialization phase for the target problem. In this way, the procedure makes use of the previously learned knowledge to boost the model’s performance and efficiency in solving the target classification problem.

The tournament selection process is utilized to determine the optimal individual within the population by randomly gathering a subset of individuals for competition based on their fitness scores. As is the case in each round, the tournament’s winner is the contestant with the highest fitness score, which is practically a measure of accuracy. The decision tree from that winner, which is essentially the best possible candidate solution, is thus designated as the best individual in the whole lot. The decision tree from the tournament’s winner is subsequently combined with the starting original population meant for the problem under consideration. The superior starter solution’s fitness advantage is therefore inherited, thereby elevating the overall effectiveness of the genetic programming algorithm at hand towards resolving the classification duly.

Over subsequent generations, the genetic programming

algorithm will shift the population from source problem-solving to target problem-solving with the help of such genetic operators as crossover and mutation that help to evolve the population, introducing new individuals while preserving the best individual of the source problem-solving. Thus, the evolutionary algorithm will take advantage of the knowledge and adaptation gained from the source problem to solve the target problem.

Nonetheless, the tournament selection mechanism will still be functioning and determining the strongest individual based on fitness score to ensure the extraction and transfer of the superior-performing population to the target problem. This initial transfer kickstarts the evolutionary process, bolstering the algorithm's efficacy in tackling the challenges specific to the target problem domain.

### 10.2. When it is Transferred:

Within the process of genetic programming, transfer learning plays a role at the beginning of solving a target problem. In particular, the collaboration of solving a problem in the target domain starts with utilising a majority of the initial population of the algorithm in the form of individuals, which are decision tree classifiers acquired from the source domain. Indeed, the goal of transferring the problem-solving abilities of the given individuals that have already demonstrated expertise in the solution methodology of the source problem is to accelerate the evolutionary procedure to tackle the target problem.

By leveraging the knowledge and insights from a source domain in genetic programming algorithms has the potential to enhance their ability to acquire beneficial adaptations and solutions, thereby expediting the convergence towards effective solutions in the target problem space. This approach capitalizes on the transfer of learned principles and adaptability from one domain to another, optimizing the algorithm's performance and solution-seeking capabilities.

The knowledge transfer process is supported in the initial phase of the genetic programming process for the target problem, characterized by the researchers as follows: population of decision trees obtained in the source problem is directly injected. This helps to use the previous experience to boost and speed up the optimization process for the new target problem and permits genetic programming to apply the knowledge and experience accumulated from the evolution in the source domain.

By leveraging the expertise and effectiveness of the best-performing individuals at the onset, the algorithm establishes a strong foundation that allows for greater efficiency and continuous improvement in addressing the complexities of the target problem. This approach not only accelerates the learning process but also promotes adaptability and performance optimization in the algorithm's problem-solving strategies, thereby contributing to enhanced overall performance and effectiveness in tackling challenging tasks.

### 10.3. How it is Transferred

Knowledge transfer in the genetic programming algorithm on the target problem requires the population to be initialised

with a set of decision tree classifiers that have been trained in advance on the source problem. Thus, the obtained classifiers serve as a background for the optimisation process to begin. In such a way, the algorithm obtains the already existing knowledge as well as the experience in solving the problem that the source domain possesses. Therefore, the evolutionary process that aims to move towards the potential effective solutions to the target problem is accelerated.

This strategic integration of pre-existing knowledge aids in kickstarting the optimization process and guiding the algorithm towards more efficient and effective solutions for the target problem. Hence, at "genetic programming with transfer learning function", individuals are transferred from the source problem to a target problem, such that at the beginning, a decision tree classifier is trained on the source problem using the Decision Tree Classifier and is then incorporated into the population as the initial individual to start the tree optimisation on the target problem. Consequently, the algorithm initialises the population with a diverse set of individuals whose chromosomes have been evolved in the past at the source problem instead of random genotypes. This way, the algorithm optimises the tree from the transferred individual, significantly decreasing the generation level as it finds the correct tree.

The population of the source problem's decision tree is cloned to generate other copies, determining the population used in initialising the target problem. Therefore, the source problem's knowledge is accumulated and integrated into the initial population. Next, the fitness evaluation function based on accuracy, precision, and recall is used to test each individual based on the target problem's training and testing data. Tournament selection is used to select the suggested individual for reproduction based on the score to promote the performance of better individuals.

Genetic variation is introduced with crossover and mutation operations to create offspring individuals, replacing some individuals in the population while maintaining a constant population size across generations. The population evolves through several generations through the iterative process of selection, reproduction, and replacement. The performance metrics of the best individual in every generation are recorded. Thus, knowledge information from the source problem to the target problem is implicitly transferred during the initialization. Individuals with inherited knowledge, the higher-performing ones on the source problem, eventually contribute to the performance enhancement of the population on the target problem.

### 10.4. Target GP Algorithm

The target Genetic Programming (GP) algorithm focuses on optimizing decision tree classifiers specifically for classification tasks. It utilizes genetic programming methodologies such as initializing a population of decision trees, assessing fitness based on metrics like accuracy, precision, recall, and confusion matrices, employing tournament selection for reproduction, and implementing crossover and mutation operations. Through multiple generations, the algorithm works to evolve the decision tree population, enhancing



their performance on the specific classification problem at hand. Each iteration refines the decision trees through genetic programming techniques to improve overall efficacy in tackling the target task.

Within the "genetic programming with transfer learning" function, the Decision Tree Classifier plays a crucial role as a key component in constructing decision tree models essential for the genetic programming procedure. Specifically, the Decision Tree Classifier is instantiated with parameters like `maxDepth`, which dictates the maximum depth allowed for the tree, thereby influencing the complexity to prevent over-fitting. Each decision tree within the population serves as a potential solution to the problem at hand, with the genetic programming approach focusing on evolving these trees to optimise performance for the given task.

Throughout the genetic programming process, every decision tree within the population is trained utilising the designated training data sets, "target X train" and "target y train". The fit method of the Decision Tree Classifier is employed with the training data, facilitating the construction of the decision tree model. This process involves iteratively dividing the feature space by using the training samples and their respective labels to establish the structure of the decision tree, paving the way for effective learning and model development within the genetic programming framework.

Following the training phase, the performance assessment of each decision tree occurs by utilising the testing data sets ("target X test" and "target y test"). The predict method of the Decision Tree Classifier is applied to make predictions on the testing data by leveraging the trained decision trees. These predicted outcomes are then matched against the actual labels to compute diverse performance metrics, including accuracy, precision, recall, and confusion matrix, offering insights into the effectiveness and predictive power of the decision tree models within the genetic programming framework.

The derived performance metrics act as the fitness scores allocated to every decision tree within the population, aiding in the evaluation of each candidate solution's proficiency in addressing the target problem. By assessing these fitness scores, the effectiveness of the decision trees becomes apparent, ultimately leading to the identification of the top-performing individual within the population. Typically, the decision tree with the highest fitness score, commonly linked to accuracy, garners recognition as the most successful solution to the given problem within the genetic programming environment.

In the context of genetic programming, the Decision Tree Classifier assumes a pivotal role as it acts as the linchpin for the creation, training, and assessment of decision tree models across the population. Leveraging its capability to acquire insights from data and provide accurate predictions, this classifier emerges as a fitting selection for tackling classification challenges prevalent in genetic programming scenarios. This underscores its significance in enhancing the algorithm's efficacy and driving successful outcomes in diverse problem-solving endeavours within the genetic

programming framework.

## 11. Experimental Setup

### 11.1. Parameter Values:

**Max Depth = 5, Population Size = 20, Tournament Size = 3, Mutation Rate = 0.1, Generations = 12, Crossover Rate = 0.8**

In decision-making, manipulating the "maxDepth" parameter in the Decision Tree Classifier enables the identification of complex data patterns; thus, one must be cautious to avoid overfitting. Generally, starting with a value between 4 and 8 is suggested. Also, over-increasing the population size in a genetic programming keeps diversity in the population and scales up exploration, which means broader parts of the solution grid space can be explored. An individual range of between 20 and 30 will work even better at maintaining the balance. Furthermore, fine-tuning the tournament size can strategically manage the balance between selection pressure and diversity, with a tournament size of 3 to 5 offering a suitable equilibrium between exploration and exploitation in the genetic programming process.

As far as the mutation rate is concerned, it is found to be more challenging to balance all the factors associated with the successful application of this characteristic. Therefore, setting a moderate mutation rate, at the value of 0.1 to 0.2, is deemed necessary to promote exploration without accelerating the inevitably fast destabilisation of the population dynamics.

Furthermore, the number of generations greater than 10, though not significant, also plays a role in enabling the algorithm to search the search space further; specifically, the values of 12 to 20 are recommended. Finally, the third parameter, the crossover rate, should be maintained between 0.6 and 0.8 to ensure genetic diversity and avoid too much focus on promising solutions at the expense of the diversity of the population.

Through the careful calibration of parameters and amplification of the generation count, the objective is to bolster the algorithm's adeptness in scouring for superior solutions across an elongated timeframe. Exploration of various parameter setups and examination of parameters' performance could help to identify which settings would provide the mileage to achieve better results. Such a process represents a methodical search for a solution to increase the algorithm's performance that aims to foster the generation of better results through experimental variation.

### Why same parameters for both the source and the target:

Ensuring that parameter values in both the source and target problems remain constant results in multiple advantages. Not only is the experimental setting unified, but the performance of the GP algorithm with or without TL can be compared, and any resultant differences are well attributed to the presence or absence of transfer learning. As a result, not only is a fair assessment possible, but implementation is also more straightforward.

It is achieved through the reduction of concerns over multiple parameter sets, enabling the reusability of code, and focusing on parameter tuning to ensure a tuned GP that improves general performance in both cases. Ultimately, the shared parameters approach reduces experimental complexity, allowing researchers to focus on evaluating the impact of transfer learning without being burdened by the complexities associated with parameter management and tuning.

## 12. experimental setup

### 12.1. Technical Specifications

- **The datasets (source and target):** are assumed to be properly formatted CSV files containing features and labels for classification tasks. Source is "diabetes.csv" file while the target is also the same "diabetes.csv" file because the problem is more related. We used the same dataset as the initial stage of our research to see how well our model can perform.
- **Programming Language:** Python environment with necessary libraries (e.g., scikit-learn, pandas, numpy, matplotlib and memory profiler).
- **Machine Specifications:** The simulations were conducted on a computer with the following specifications:
  - **Processor:** Intel Core i5
  - **Memory:** 8 GB RAM
  - **Operating System:** Any OS capable of running Python and required libraries. But for this project, the code is ran in anaconda (Jupyter node book) on windows 11.
  - **Computational Resources:** Adequate CPU and memory resources to handle the GP algorithm's computational demands.

#### 12.1.1. Software and Libraries

- **Data Preprocessing:** Data preprocessing tasks, such as reading CSV files and splitting the dataset, were performed using the pandas library.
- **Machine Learning:** The scikit-learn library was utilized for implementing the decision tree classifier, train-test split, and evaluation metrics (accuracy, precision, recall).
- **Genetic Programming Implementation:** The genetic programming algorithm was implemented using custom Python code, leveraging the functionalities of numpy for array manipulation and matplotlib for plotting.

#### 12.1.2. Experimental Procedure

- The dataset (name: diabetes.csv) was preprocessed, with features and labels extracted for input into the genetic programming algorithm.
- The algorithm was executed with the specified parameter values, including population size, maximum depth, tournament size, mutation rate, and number of generations.
- Performance metrics (accuracy, precision, recall) were computed for each generation to evaluate the effectiveness of the algorithm in solving the classification task.
- Confusion matrices were generated to analyze the classification results and understand the model's predictive behavior.

- The runtime of the algorithm, as well as the time taken for each generation, was recorded to assess computational efficiency.

## 13. Runtime

### 13.1. Runtime Measurement:

Runtime is measured using the time module. Specifically, the time.time() function is used to record the current time at the beginning and end of a section of code, allowing for the calculation of elapsed time.

In the code, the start time is recorded before the execution of a specific task (e.g., evaluating fitness, generating a new generation), and the end time is recorded after the task is completed. The difference between the start and end times gives the elapsed runtime for that task.

The total runtime for the entire process is obtained by summing up the individual task runtimes.

### 13.2. Memory Usage Measurement:

Memory usage is measured using the memory usage function from the memory profiler module. This function is used to monitor the memory usage of a specific section of code or function.

Within the code, the memory usage function is called to record the maximum memory usage during the execution of a particular task. By wrapping the code of interest with the memory usage function, the memory usage at various points in the code can be monitored.

The maximum memory usage observed across all monitored points is typically considered as the total memory usage for the task.

### 13.3. Implementation:

In the code provided, runtime and memory usage are measured within loops corresponding to different generations of a genetic programming algorithm. At the beginning of each loop iteration, the start time is recorded, and memory usage is monitored using the memory usage function. After the completion of the specific task (e.g., fitness evaluation, generation update), the end time is recorded, and the memory usage is again monitored. The difference between start and end times gives the runtime for the task, and the maximum memory usage observed during the task is recorded as memory usage. These values are accumulated across all loop iterations to obtain the total runtime and memory usage for the entire process.

### 13.4. Purpose

Runtime and memory usage measurements provide insights into the performance and resource requirements of the genetic programming algorithm. Monitoring runtime helps in understanding how long each task or iteration takes to execute, which can be crucial for optimizing the algorithm's efficiency. Tracking memory usage is essential for identifying potential memory leaks or inefficiencies in the algorithm's memory management, particularly when dealing with large datasets or

complex models. Overall, runtime and memory measurements aid in evaluating the scalability, performance, and resource utilization of the genetic programming algorithm.

## 14. Results and Analysis

We conducted a comprehensive experiment to assess the performance of our model with and without transfer learning. Specifically, we executed a loop of 12 runs, each consisting of 12 generations, and recorded various metrics for each generation, including overall accuracy, overall precision, total runtime, total memory usage, precision, recall, F1-score, and support. However, for the purpose of this report, we have primarily focused on the accuracy and precision metrics, as these are critical indicators of the model's effectiveness. The results obtained from each run are presented below.

Run	Accuracy without TF	Accuracy with TF	Precision without TF	Precision with TF
1	0.79	0.79	0.75	0.75
2	0.79	0.79	0.75	0.75
3	0.80	0.80	0.75	0.75
4	0.79	0.80	0.75	0.75
5	0.79	0.79	0.75	0.75
6	0.79	0.80	0.75	0.75
7	0.79	0.80	0.75	0.75
8	0.80	0.79	0.75	0.75
9	0.79	0.80	0.75	0.75
10	0.79	0.80	0.75	0.76
11	0.79	0.80	0.75	0.75
12	0.79	0.79	0.75	0.75
Average	0.79	0.80	0.75	0.75
Best:				
10	0.79	0.80	0.75	0.76

Based on the results obtained from our experiment, a comparison was conducted between the accuracy of our model with and without transfer learning. Notably, the model without transfer learning achieved an accuracy of 80 percent in only two runs, while the remaining 10 runs resulted in an accuracy of 79 percent. Furthermore, the precision of this model remained constant across all generations, at a rate of 75 percent. In contrast, the model with transfer learning demonstrated a more varied performance, with seven runs achieving an accuracy of 80 percent and five runs achieving an accuracy of 79 percent. Interestingly, the model with transfer learning also exhibited a more nuanced precision profile, with only one run resulting in a precision of 76 percent.

Notably, run number 10 exhibited a superior performance compared to the other runs. Specifically, the model with transfer learning in this run achieved an accuracy of 80 percent and a precision of 76 percent, which is the highest precision observed in all runs. This exceptional performance suggests that the model was able to effectively leverage the knowledge gained from the source population to improve its performance on the target population, leading to a remarkable improvement in both accuracy and precision.

## 15. Runtime

Totals	Without Transfer	With Transfer
Runtime	39.33 seconds	2.56 seconds
Memory	53030.72 MB	51968.48 MB

During each run, we initialised time and space measurements to track the processing time and memory usage of the genetic algorithm and knowledge transfer process. Our analysis revealed that the model without transfer learning required a significantly longer runtime, taking approximately 39.33 seconds to complete. Additionally, it utilised a substantial amount of memory, reaching a peak of 53,030.72 MB. In contrast, the model with transfer learning demonstrated a remarkable reduction in runtime, completing in just 2.56 seconds.

Although it also consumed a considerable amount of memory, reaching a peak of 51,968.48 MB, our findings indicate that our model's ability to transfer knowledge in a fast and efficient manner was a significant contributing factor to its improved performance.

### 15.1. Conclusion

In conclusion, the experiment validates the efficacy of transfer learning in enhancing the model's performance. The model with transfer learning had better prediction accuracy and precision levels than the one without transfer learning. These outcomes signify that the acquired knowledge from the source population was successfully transferred to the target population. The usefulness of the knowledge transfer process was more evident in the substantial reduction in both runtime and memory usage.

In aggregate, the outcomes imply that transfer learning is a viable approach to enhancing the machine learning model's performance, especially in datasets with limited datasets.